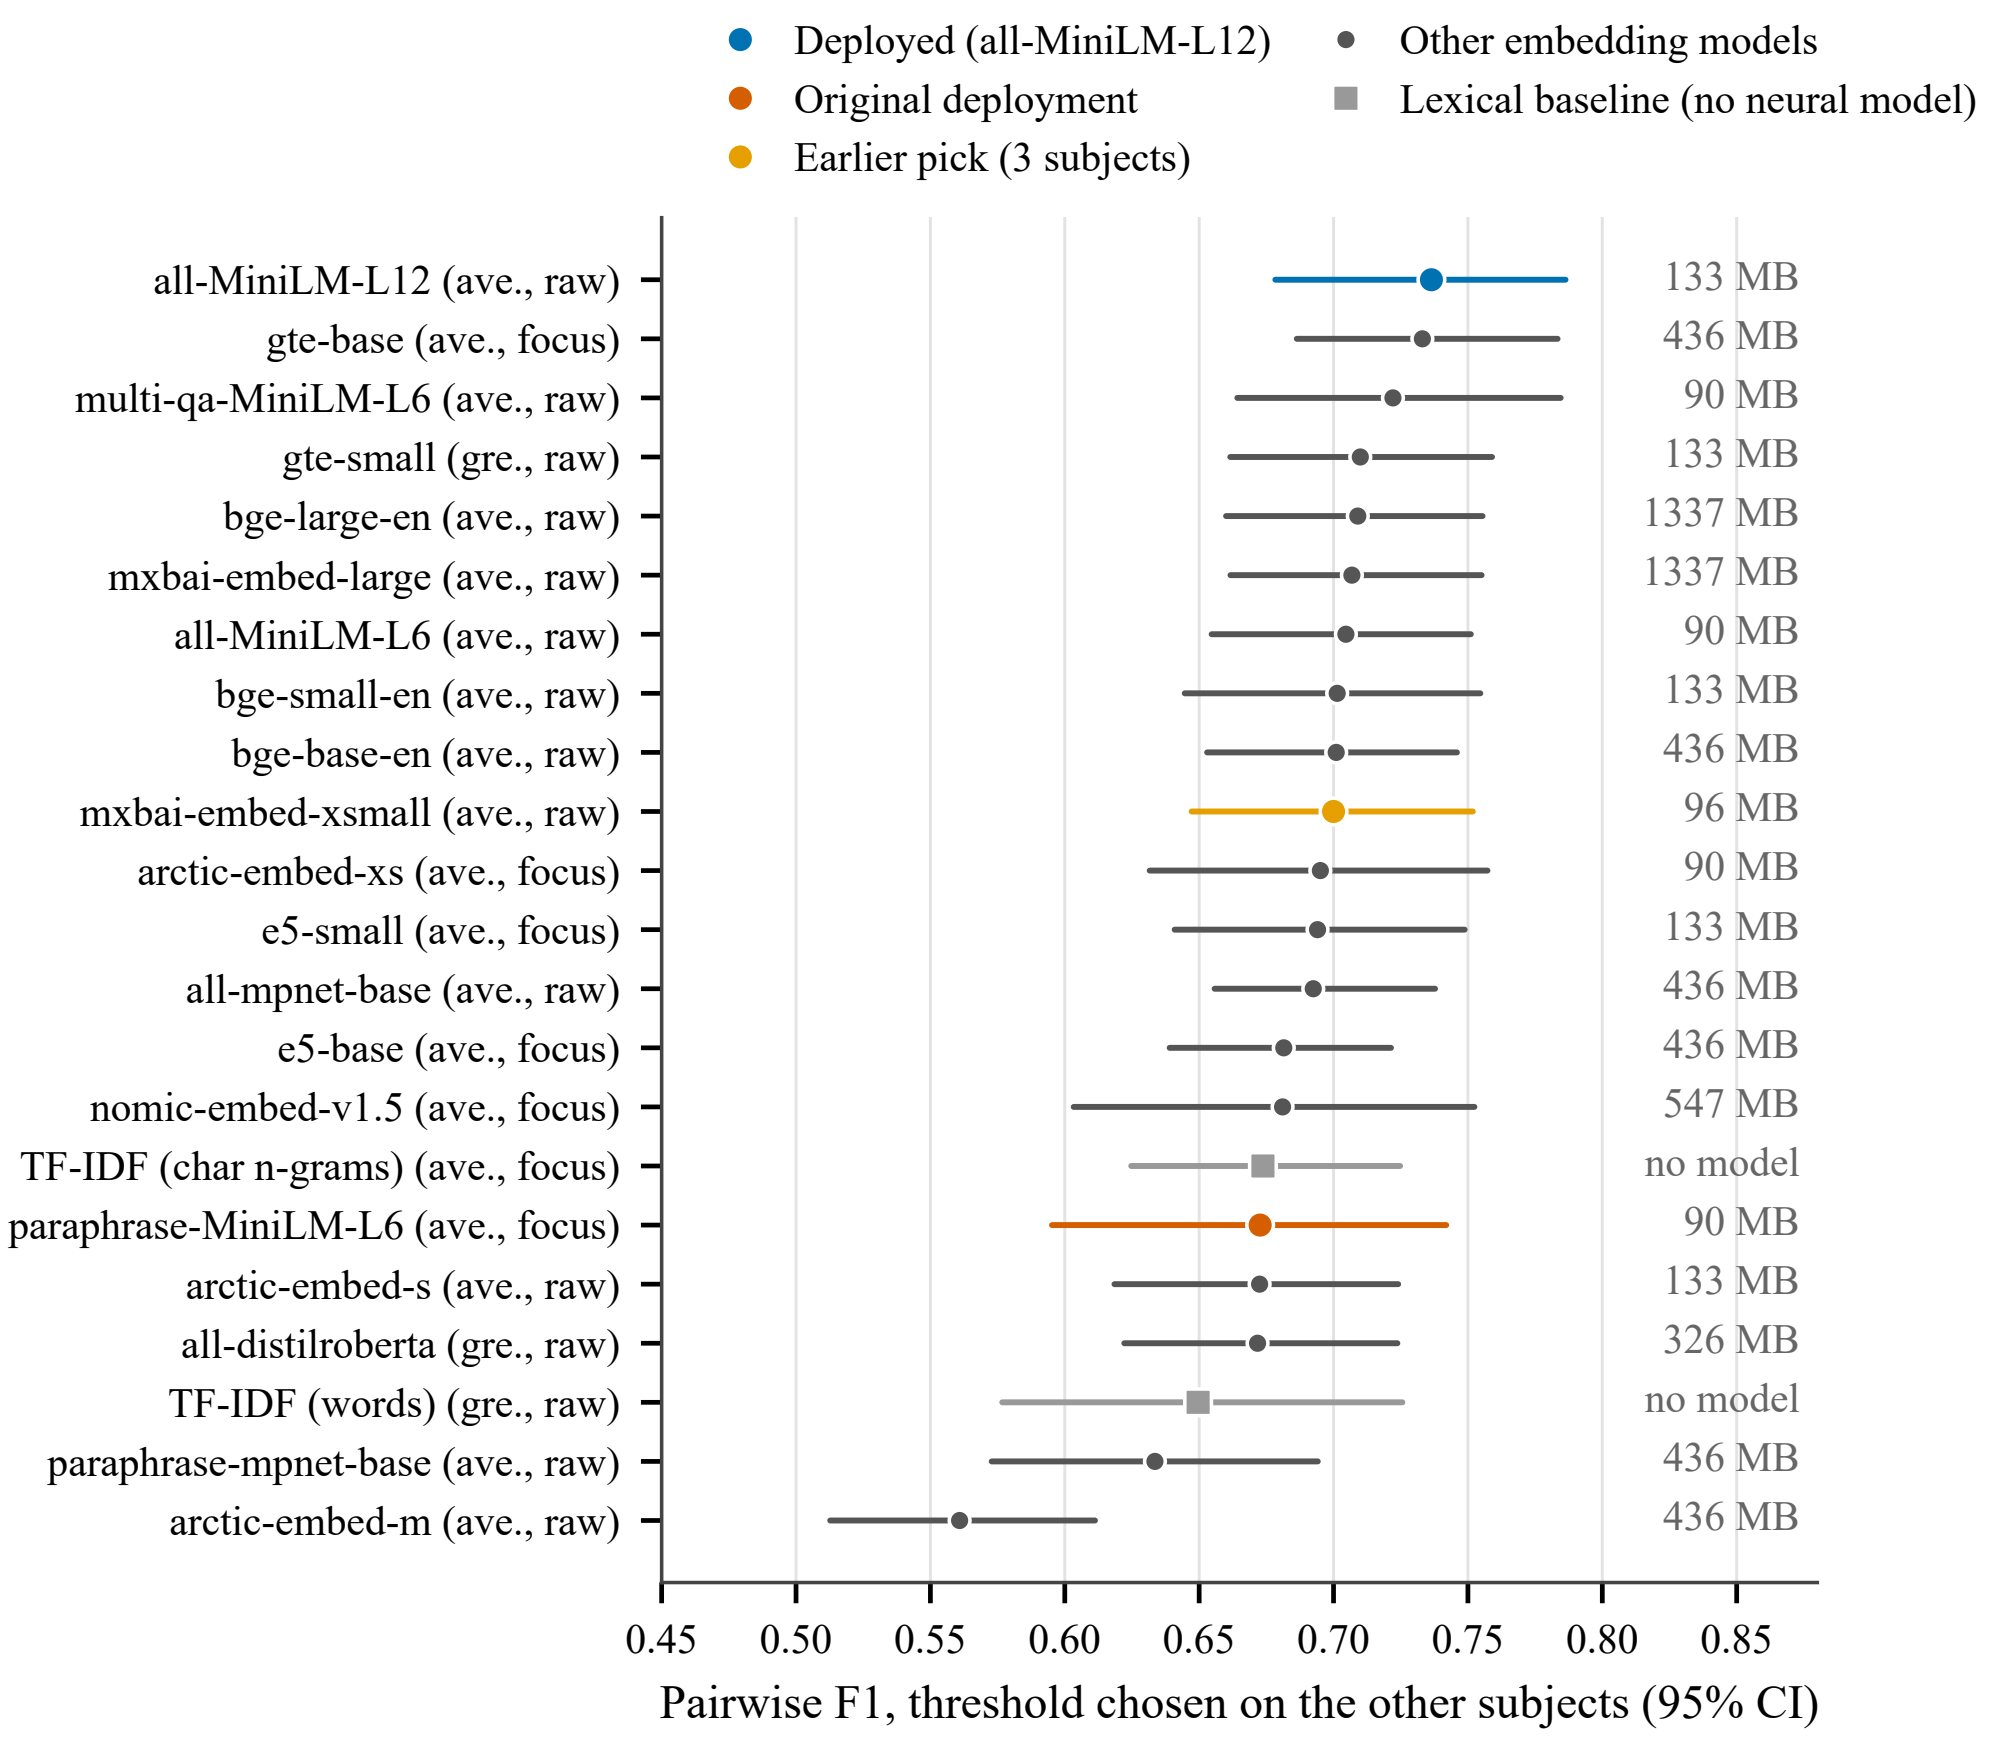
Deployed (all-MiniLM-L12)
Original deployment
Earlier pick (3 subjects)
Other embedding models
Lexical baseline (no neural model)
all-MiniLM-L12 (ave., raw) 133 MB
gte-base (ave., focus) 436 MB
multi-qa-MiniLM-L6 (ave., raw) 90 MB
gte-small (gre., raw) 133 MB
bge-large-en (ave., raw) 1337 MB
mxbai-embed-large (ave., raw) 1337 MB
all-MiniLM-L6 (ave., raw) 90 MB
bge-small-en (ave., raw) 133 MB
bge-base-en (ave., raw) 436 MB
mxbai-embed-xsmall (ave., raw) 96 MB
arctic-embed-xs (ave., focus) 90 MB
e5-small (ave., focus) 133 MB
all-mpnet-base (ave., raw) 436 MB
e5-base (ave., focus) 436 MB
nomic-embed-v1.5 (ave., focus) 547 MB
TF-IDF (char n-grams) (ave., focus) no model
paraphrase-MiniLM-L6 (ave., focus) 90 MB
arctic-embed-s (ave., raw) 133 MB
all-distilroberta (gre., raw) 326 MB
TF-IDF (words) (gre., raw) no model
paraphrase-mpnet-base (ave., raw) 436 MB
arctic-embed-m (ave., raw) 436 MB
0.45
0.50
0.55
0.60
0.65
0.70
0.75
0.80
0.85
Pairwise F1, threshold chosen on the other subjects (95% CI)